

Echo - Product Brief

Product context

Echo is an observability copilot for HighLevel Voice AI agents. It helps teams understand whether calls are healthy, whether an agent followed its goals, and what should improve next.

Who it is for

- Operations and QA teams reviewing many calls without reading every transcript manually.
- Conversation designers checking whether an agent follows its script or business goal.
- Team leads turning repeated failures into concrete coaching or prompt changes.


The problem

Voice AI teams collect transcripts but still have to decide what went wrong, whether it matters, and how to improve the agent. Echo makes that decision visible and auditable instead of leaving it inside raw logs.

Product objective

Close a simple Monitor -> Analyze -> Act loop: collect calls, evaluate configured goals, show evidence, and turn repeated failures into a reviewable recommendation.

Main workflow



Core feature set

Feature	User value
Call ingestion	Completed calls arrive from HighLevel webhooks or call-log backfill.
Baseline monitoring	Every call gets completion, sentiment, talk share, turns, and other call signals.
Agent settings	A user defines manual pass/fail criteria in one compact row-and-modal workflow per agent.
Evidence-linked evaluation	Each criterion shows its result, rationale, score, and supporting transcript turns.
Review actions	Call segments become follow-up, training, or compliance work that can be resolved or dismissed.
Recommendations	Recurring failures become ranked prompt or script suggestions linked to evidence calls.

Product experience

- Overview answers: are calls healthy, and where should I look first?
- Agent settings answers: what does success mean for this agent?
- Calls answers: which calls need attention?
- Call detail answers: what happened, and can I verify the finding?
- Agent detail answers: what failure pattern should I improve?

Product decisions

- Criteria are manual and per-agent, so Echo never invents custom verdict rules.
- Baseline monitoring stays visible when criteria are empty; custom pass/fail appears only after setup.
- Evidence links every important result back to the transcript.
- Recommendations are copyable suggestions and never silently change a live agent.

AI-first, not AI-only

AI handles language: quality interpretation, ambiguous criterion judgements, evidence rationales, and recommendation wording. Deterministic code handles ingestion, metrics, thresholds, weighted scoring, validation, and aggregation so the numbers remain reproducible.

Install in HighLevel

- Create or install the Echo marketplace app for the HighLevel sandbox.
- Add Echo as a Custom Page or Custom Menu Link inside the account.
- Authorize the location and configure the completed-call webhook. Echo then syncs agents and collects calls.

MVP boundaries

Seeded calls support demonstration, while HighLevel ingestion and LLM analysis are the live product paths. Prompt write-back is not automated. A lightweight text tag can be added later for filtering and grouping. Pre-computation tables are also a future optimization; they are not included yet because direct aggregation keeps the MVP simpler.

Tech stack

Layer	Technology
Frontend	Vue 3, Vite, Tailwind CSS, TanStack Query
Backend	Node.js, Fastify, Drizzle, Inngest, Zod
AI	Structured LLM evaluation and recommendation generation
Data	PostgreSQL via Supabase

The assignment brief is preserved separately in docs/assignment-brief.pdf.